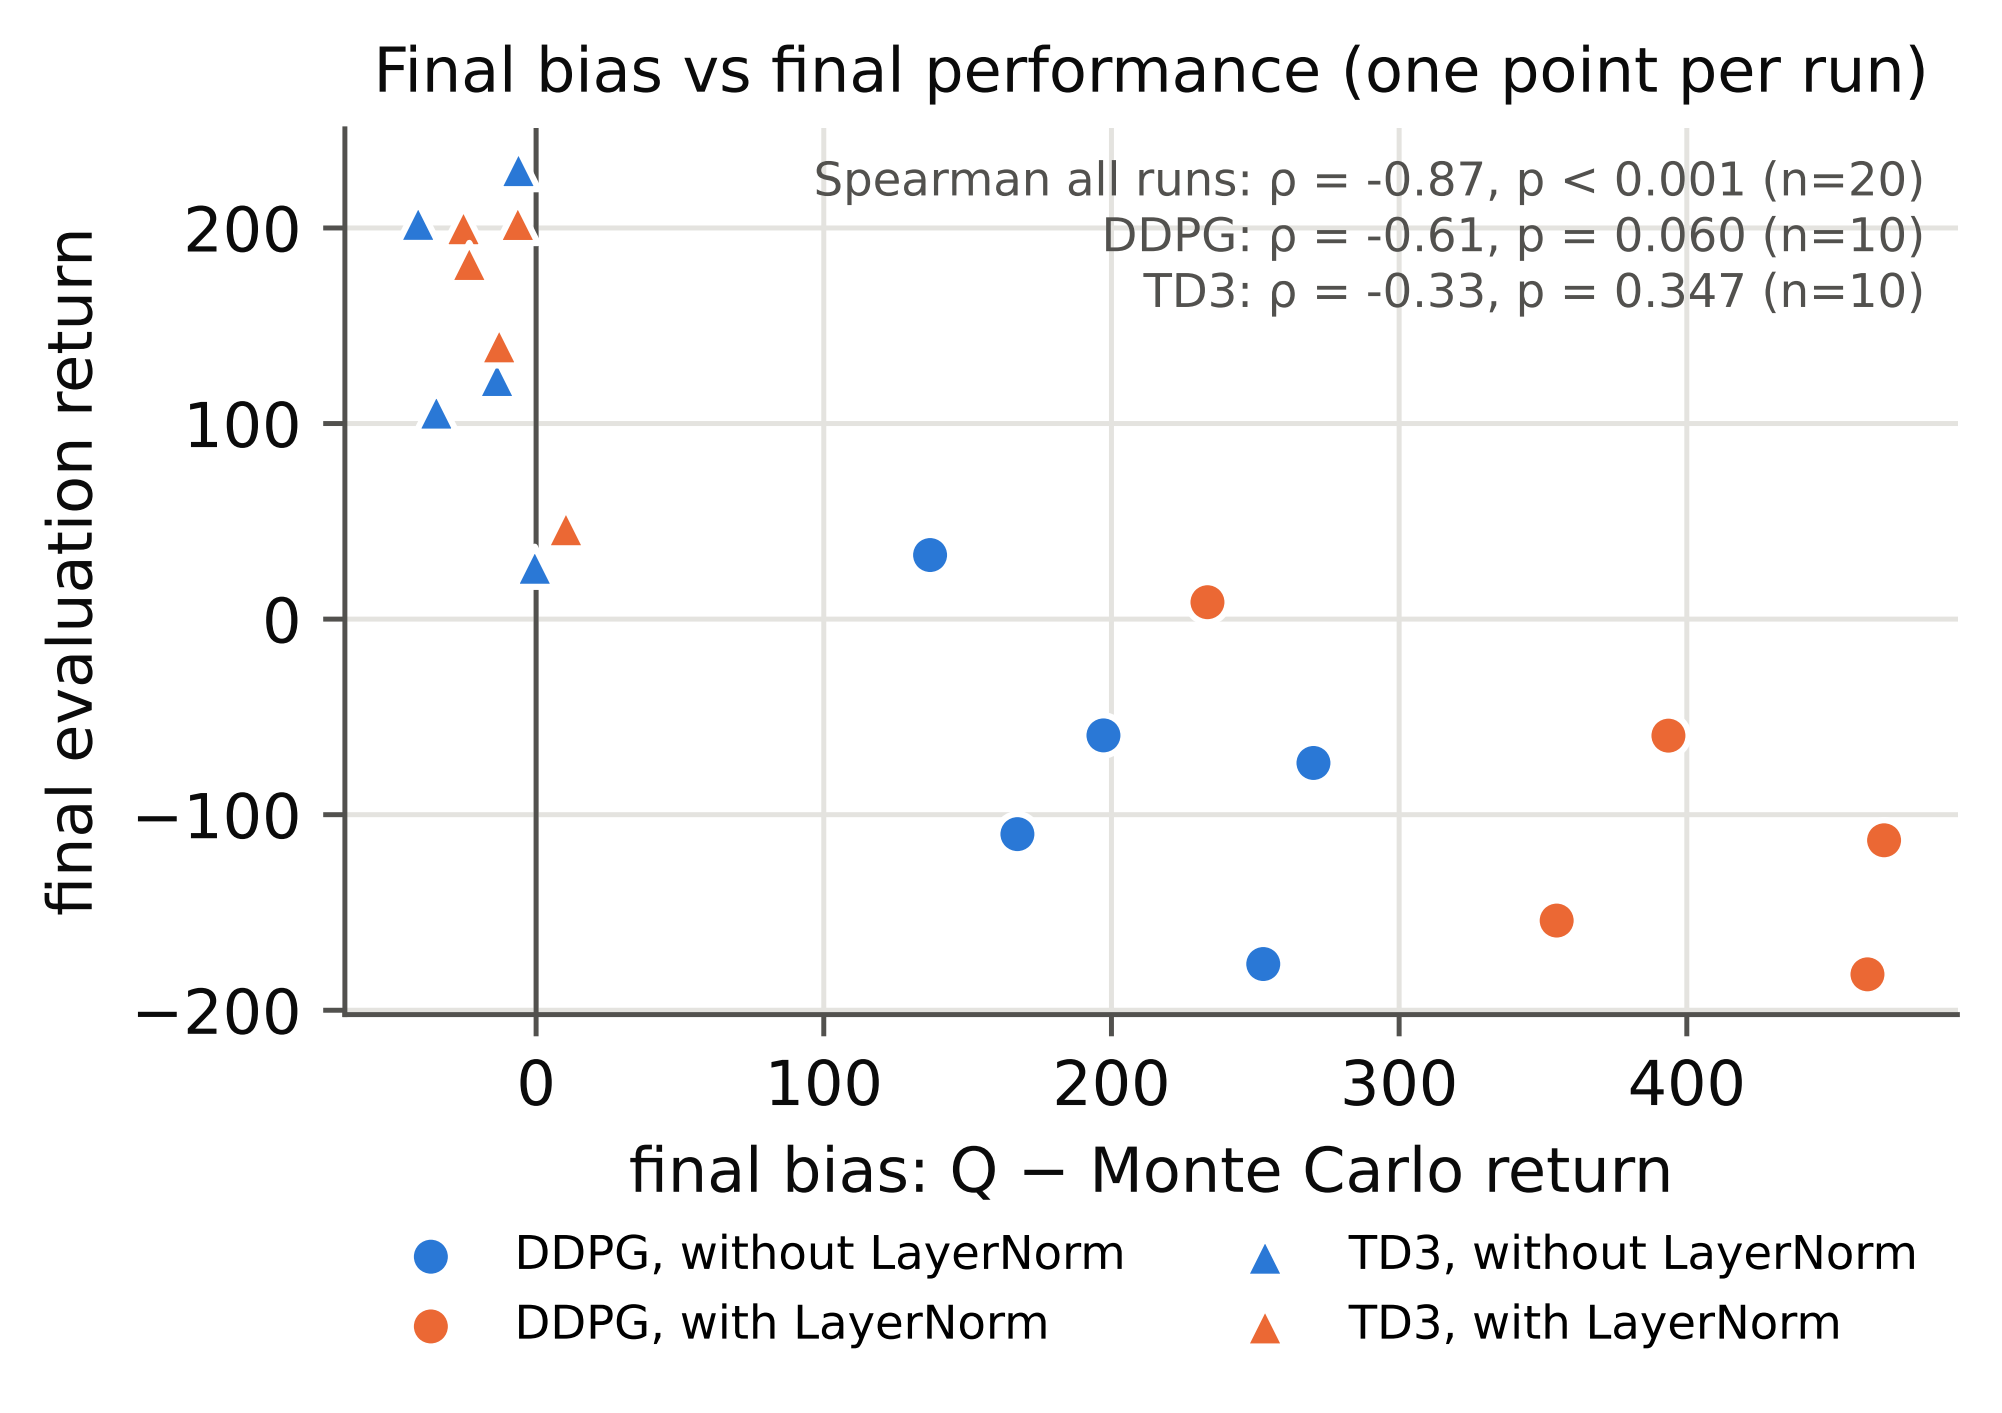
Final bias vs final performance (one point per run)
Spearman all runs: ρ = -0.87, p < 0.001 (n=20)
DDPG: ρ = -0.61, p = 0.060 (n=10)
TD3: ρ = -0.33, p = 0.347 (n=10)
final evaluation return
200
100
0
−100
−200
0
100
200
300
400
final bias: Q − Monte Carlo return
DDPG, without LayerNorm
DDPG, with LayerNorm
TD3, without LayerNorm
TD3, with LayerNorm