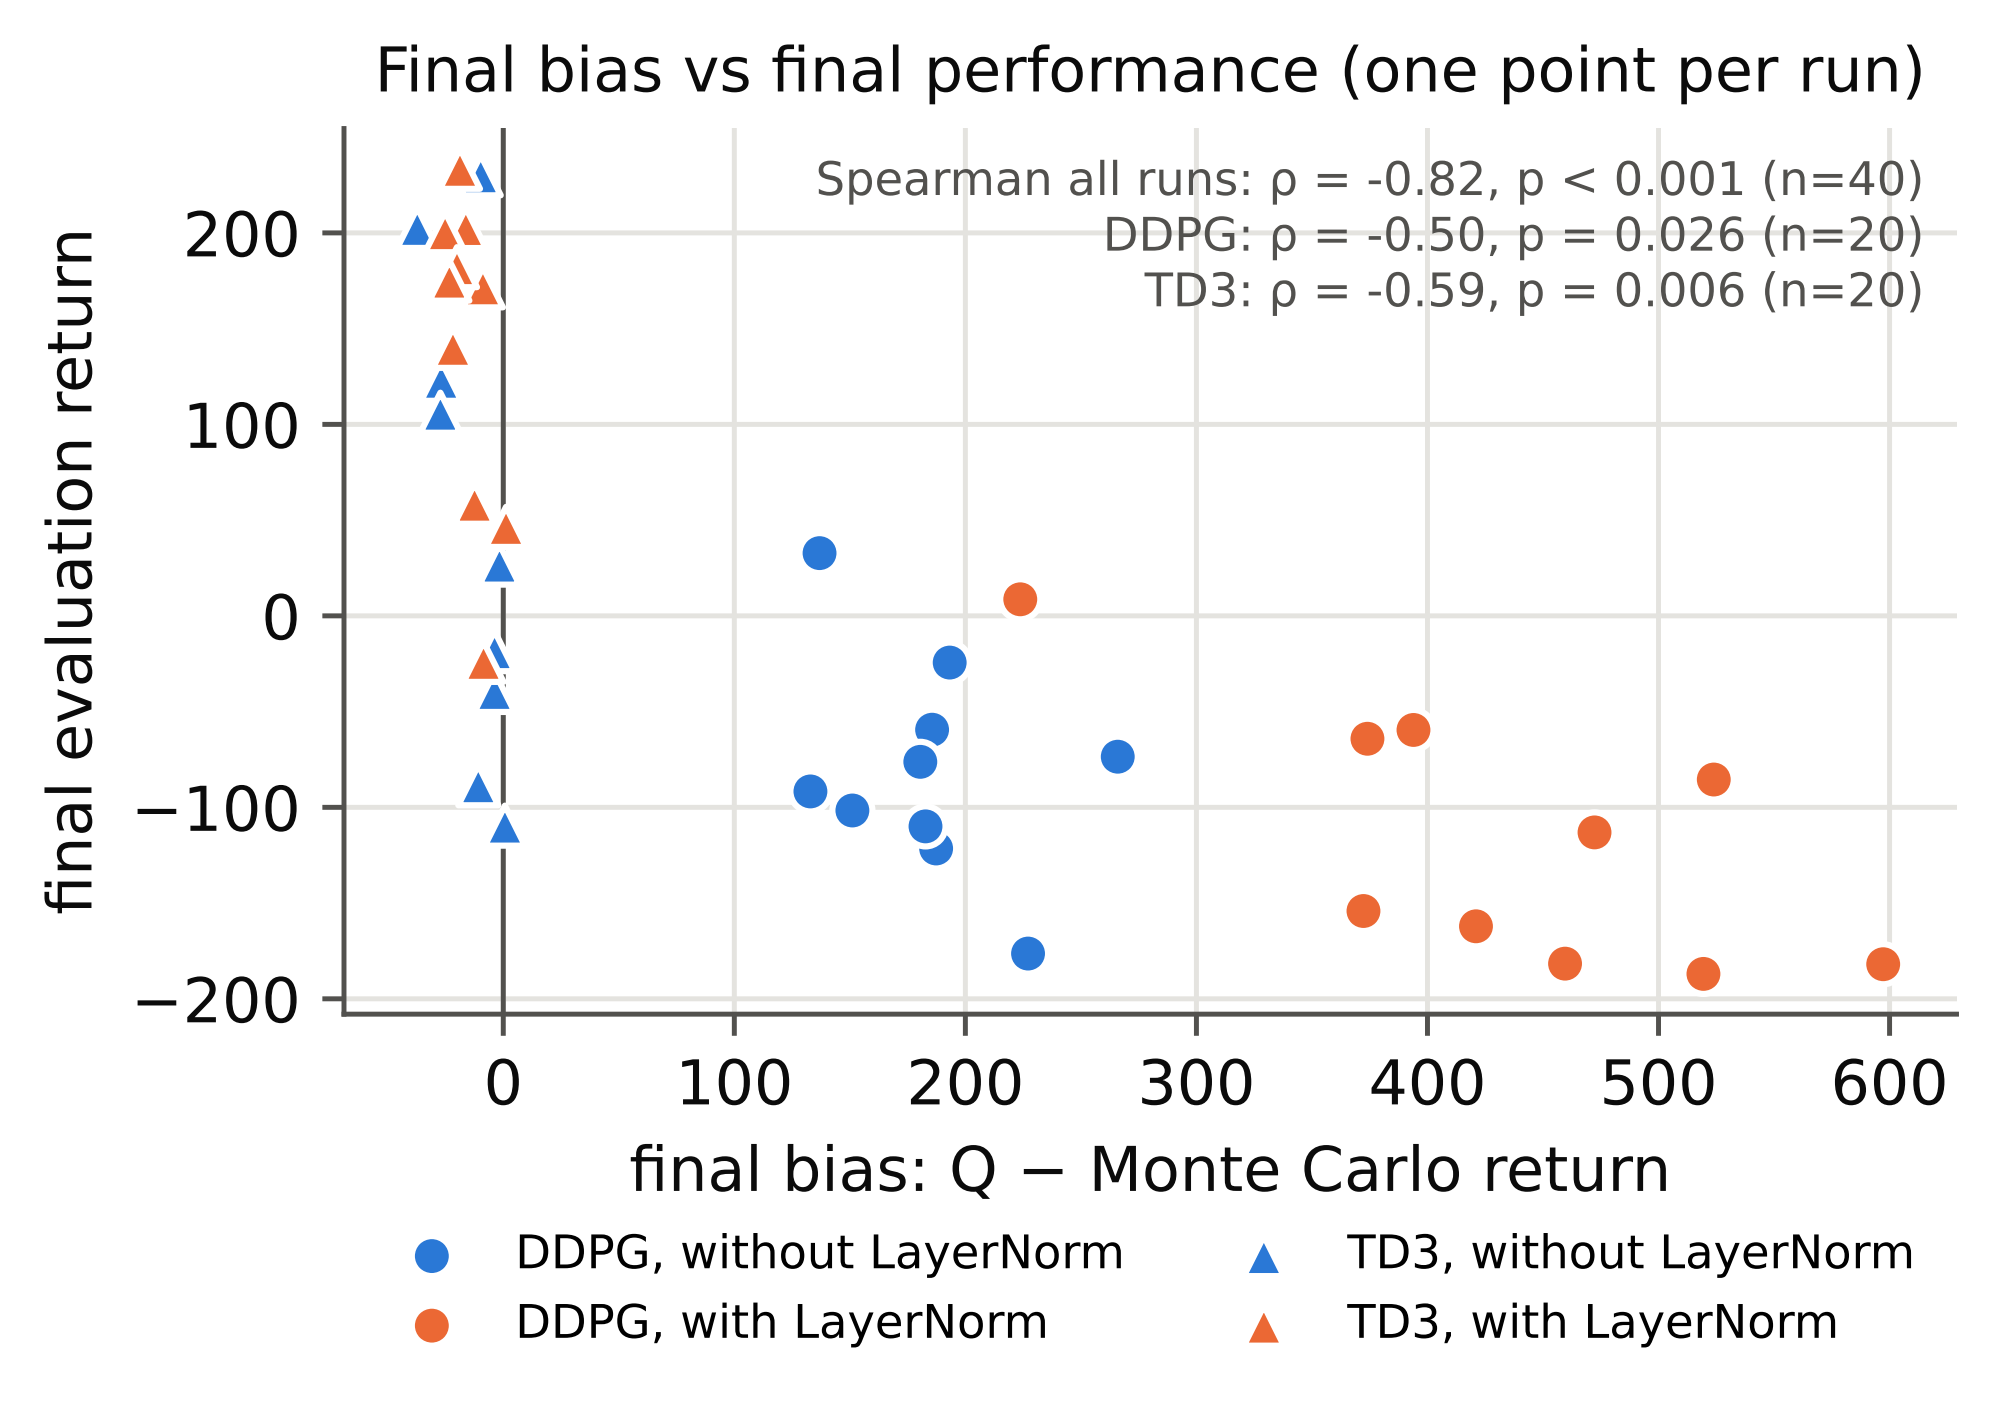
Final bias vs final performance (one point per run)
Spearman all runs: ρ = -0.82, p < 0.001 (n=40)
DDPG: ρ = -0.50, p = 0.026 (n=20)
TD3: ρ = -0.59, p = 0.006 (n=20)
final evaluation return
200
100
0
−100
−200
0
100
200
300
400
500
600
final bias: Q − Monte Carlo return
DDPG, without LayerNorm
DDPG, with LayerNorm
TD3, without LayerNorm
TD3, with LayerNorm